

Changelog

All notable changes to this project will be documented in this file.

The format is based on [Semantic Versioning](#).

[1.1.0] - 2025-08-10

Changed

- Added **version** field to **GET /health** response (API versioning introduced).
- Promoted model endpoints from **/online/model/*** to root level: **GET /model/info**, **POST /model/export**.
- Replaced **POST /online/training/restart_policy** with **PATCH /online/training/restart-policy** for RESTful semantics.
- Prediction endpoints (**POST /predict**, **POST /recommend**) now return HTTP 503 with **status: model_unavailable** when the model is not loaded instead of returning random fallback scores.

Removed

- Deprecated **POST /online/model/retrain** endpoint (non-online full retrain placeholder) removed.
- Removed embedded changelog section from **README.md** (now separate file).

Added

- Root-level model export (**POST /model/export**) and aggregated model info now include incremental trainer state when active.
- Separate **CHANGELOG.md** file.
- Stronger path sanitization for model/config directories.

Migration Notes

- Update any clients calling **/online/model/export** or **/online/model/info** to use root-level equivalents.
- Remove usage of **/online/model/retrain**; implement offline retraining workflow separately if needed.
- Adjust clients expecting restart policy change via POST to use PATCH at new hyphenated path.
- Adjust client error handling: treat 503 **model_unavailable** as a retrievable state and avoid using stale or placeholder predictions.

[1.0.0] - 2025-08-01

Added

- Initial release with core recommendation endpoints: **POST /recommend**, **POST /predict**.

- Health endpoint `GET /health` (without version field).
- Online learning endpoints: data ingestion, incremental training start/stop/status/logs.
- Streaming (Kafka) utilities and incremental updates listing.
- `/online/model/export` and `/online/model/info` (now relocated in 1.1.0).